

The Major AI Ecosystems

D-Lab AI Pulse, Fall 2026: Session 2

Materials: https://github.com/dlab-berkeley/AI_Pulse
Workshop suggestions: <https://forms.gle/x8KfejgCViE7t6iQA>

15 September 2026

Where Session 1 Left Off

- AI has developed well past the chat box: agents, projects, and connectors
- Getting the most out of any of it takes an application, and a subscription
- Which raises three questions:
 - What is the difference between the environments
 - What is the difference between the plans
 - Is it worth paying more, given what I actually do
- Today is about answering all three

Today's Roadmap

- We look at the two main environments: OpenAI and Anthropic
- Part 1, the basics, what each company sells and what it costs
- Part 2, what they have in common, which is most of it
- Part 3, where they differ in practice, run side by side
- Part 4, what you are agreeing to, and what happens when it changes
- Throughout, one task carried end to end: how do I keep up with AI

Part 1

The Basics

Why These Two

- Berkeley gives you Gemini and Gemini Notebook already, so that decision is made for you
- They are close to the frontier at all points
- Both build for general use, not for one narrow job, so what you learn transfers
- Both have a large community, which is where the workflows actually come from
- And both are established enough to be, somewhat, trustworthy with your work
- There are others for general use: Grok, Kimi, Mistral, DeepSeek

Two Model Lineups

- Anthropic keeps parallel tiers. OpenAI ships variants of one model

	Anthropic	OpenAI
Frontier	Fable 5.1	GPT-6 Astra
Default	Opus 5	GPT-5.6 Sol
Balanced	Sonnet 5	GPT-5.6 Terra
Cheap	Haiku 4.5	GPT-5.6 Luna

- How hard the model thinks is now a slider, not a separate family of models

What You Get, Not What You Pay

	Anthropic	OpenAI
Free	Chat, skills, 5 projects	Chat, Codex, Work, projects
\$8	none	+ more files, 3 scheduled jobs
\$20	+ Claude Code, Cowork, projects	+ 5 scheduled jobs
\$100	+ five times the \$20 limits	+ five times, 15 scheduled jobs
\$200	+ twenty times the \$20 limits	+ twenty times

- OpenAI has one rung Anthropic does not: an \$8 tier, which may carry ads

Two Different Vocabularies for Usage

- Anthropic uses one vocabulary: a rolling five-hour session, plus weekly caps
 - Chat, Claude Code and the desktop app all draw from the same pool
- OpenAI uses four at once: unlimited per model, credits per message, credits per million tokens, and rolling windows
 - A single deep research run costs 50 credits; an image costs 5
- On the top Anthropic model, half your weekly allowance is the ceiling
 - On the \$20 plan it is not included at all: it bills as you go
- This is the most confusing part of either product, and the hardest to compare
- When you run out you can buy more, by credits or through the API

What You Might Not Have to Pay For

Program	Who	What you get
OpenAI, Academic Researchers	Faculty or postdoc, with a preprint in the last three years	Twelve months free at the top limits, five seats
OpenAI, Researcher Access	Any researcher, reviewed quarterly	Up to \$1,000 in API credits
Anthropic, Team Plan for Scientists	Principal investigators in science and related fields	Free seats for twelve months, \$15 for the larger ones
Anthropic, AI for Science	Academic and nonprofit projects	Up to \$50,000 in credits per project
Anthropic, Life Sciences	Vetted life scientists, US institutions	The restricted model, and looser safeguards

- None of these is a student route: you get in through a supervisor's group
- Both verify against your institutional email, and both take a week or two

Part 2

What They Share

Both of Them Already Do All of This

None of it is a reason to pick one over the other.

- Projects, memory, deep research, and code execution inside chat
- Connectors to your other tools, and MCP support on both sides
- Computer use, scheduled tasks, and a million tokens of context
- Batch and caching discounts, Office integrations, and a unified desktop app

So I will show each of these once, and it will look the same wherever you do it.

Asking a Question

- Both will search the web and come back with something about as good as the other
- They cite what they used, and you can click through to the source
- You can ask follow-ups against what it found, rather than starting again
- Everyone already does this, and it is where the two are hardest to tell apart

Keeping Sources in One Place

- A Project is a folder with a memory: files, instructions, and its own history
- Papers, links and notes go in once, and every later conversation can see them
- Standing instructions live there too, so you stop retyping what you care about
- If you take one thing away today it is this, and both of them have it

Projects: The One Real Difference

- Anthropic caps the number of projects on Free and removes the cap when you pay
- OpenAI gives unlimited projects on every plan, and caps the files inside each
- Either way a project fills up. Anthropic's then starts searching your files instead of reading them all, which holds about ten times as much
- So a project is not a bottomless folder. Past a point it retrieves rather than reads, and you will not be told when that happened
- OpenAI lets you choose whether a project can see your general memory, or only its own, which matters if unpublished work lives in one and not another

What It Remembers About You

- Both keep a memory across conversations, not only inside one
- Anthropic now writes topics down as you talk rather than summarizing afterward, and shows them to you as files you can edit or delete
- It spans the chat and the agent, so something you said in one turns up in the other
- Health, race, religion, politics and gender identity are left out unless you opt in
- On by default for individuals, and off until an administrator turns it on for a team

Convergent Evolution

- Both companies built a chat app, a coding agent, and an agent for everything that is not code
- Both then merged all three into a single desktop application
- Anthropic's has Chat, Cowork and Code tabs. OpenAI's has Chat, Work and Codex
- They shipped six weeks apart, without copying each other's design

Part 3

Where They Differ in Practice

Artifacts against Canvas

- Both let the model produce a document you can edit beside the conversation
- OpenAI narrowed Canvas to writing and code in May
- Anthropic's Artifacts stayed broad: working pages you can publish and share
- An artifact that calls the model bills the person viewing it, not the person who wrote it

A Canvas, Not a Conversation

- Anthropic has a separate surface for designing things: chat on one side, a canvas on the other
- You describe what you want, then refine it by editing the canvas directly
- It exports to PowerPoint and PDF, which is the part that matters for this room
- It is still in beta, and it is switched off by default on institutional accounts
- Editing at the same time as somebody else does not reliably work yet

Turning Sources Into Something Readable

- From a folder of papers, it can write a summary of what changed and why
- You can set the shape: length, audience, what to leave out
- The same prompt in the two environments gives noticeably different writing
- This is the first point where which one you use changes what you get

Making a File

- Anthropic writes Word, Excel, PowerPoint and PDF files on every plan, including the free one
- OpenAI writes to Google Docs and Sheets, and has an Excel add-in
- PowerPoint is missing from OpenAI's desktop flow, and PDF output is undocumented
- If you send colleagues attachments, this matters more than it sounds like it should

Producing a File Someone Else Can Open

- It can turn what you have into a document, a slide deck or a spreadsheet
- One of them makes the file for you, the other sends you to Google
- Where the file lands, and who can open it, differs between them

Two Coding Agents

- Claude Code and Codex do the same job: they read your files, write, run and fix code
- Both run in a terminal, in your editor, on the web, and on a phone
- Codex is on OpenAI's free tier. Claude Code needs a paid plan
- You do not have to write code to use either. Most research use is data wrangling and file handling

Two Agents for Everything That Is Not Code

- Cwork and ChatGPT Work are the same idea aimed at documents rather than repositories
- You describe a job, it works through your files, and you review what came back
- ChatGPT Work is generally available and on the free tier
- Cwork is generally available on desktop and still in beta on web and mobile

Sending It Off to Work Alone

- You can give it a job that takes minutes rather than seconds, and walk away
- It decides what to open, what to skip, and when it has enough
- It reports back with what it found and where it came from
- This is where the two pull apart most

Connecting Your Own Tools

- Both can reach your mail, calendar, drive and notes through connectors
- Both use MCP, the same open standard, which Anthropic wrote and OpenAI adopted
- That is unusual. Competitors normally ship a rival rather than adopt yours
- Enabling any connector widens what the model can see considerably

Letting It Read Your Own Material

- Connectors let it read your mail, calendar, drive and notes
- You choose the scope before it runs: one label, one folder, one calendar
- Newsletters are the case where this genuinely saves an hour a week
- It searches rather than reads everything, so ask what it searched

Work That Runs Without You

- Both can run a job on a schedule, at most once an hour
- OpenAI caps how many you can have: three on Go, five on Plus, fifteen on Pro
- Anthropic does not cap the count, and can also trigger on a web request or a GitHub event
- So one is a scheduler and the other is closer to a small automation system

Two Kinds of Repeating Job

- One kind repeats on a clock: check this every Monday, or every five minutes
- The other repeats until a condition is true, with a limit on how many attempts it gets
- The second is the useful one, because you supply the test instead of the model judging its own work
- Anthropic's phrasing: this way it does not have to decide what is good enough
- Match the interval to how often the thing actually changes, or you pay thirty times over for the same answer

Having It Run Without You

- A job that works once can be set to repeat on a schedule
- It runs whether or not you are at the machine, and leaves the result waiting
- Look at my calendar for tomorrow and tell me what I need to prepare
- Check my inbox for anything that actually needs answering today
- Watch arXiv for new papers in my area and summarize what is worth reading

Packaging a Repeatable Job

- Once a job works, both let you save it so it repeats without rebuilding
- Anthropic calls these Skills: a folder with instructions, zipped and uploaded
- You can also record one: do the job once while it watches, and it writes the skill for you
- No code required, and it is an open standard other tools can read
- OpenAI's equivalent is Custom GPTs, and its own Skills feature is elsewhere

Where an Instruction Should Live

Put it in	Loaded	Use it for
The prompt	This turn only	A one-off request
Custom instructions	Every turn, always	Tone and format you always want
Project instructions	Inside that project	Conventions for one body of work
A saved skill	When it is relevant	A procedure you repeat

- The common mistake is keeping everything in the prompt and retyping it forever
- The other one is putting everything in the always-loaded layer, which you pay for on every single turn

The Asymmetry at the End of the Demo

- Skills are on every Anthropic plan, including the free one
- On OpenAI they are Business and Enterprise only
- So the last stage of today's task can be packaged on one side and not the other
- And not because of price: there is no personal plan that unlocks it

Part 4

What It Means for You

- On a personal paid or free account, Anthropic asks first. Training is off unless you turned it on
- OpenAI's is on by default. You have to go and turn it off
- If you leave it on, your conversations can be retained for up to five years
- On Team, Enterprise, Education and API plans, neither company trains on your content
- Both have a private mode that is never trained on: Incognito, and Temporary Chats

Turning Training Off Is Not Deleting Anything

- These are separate controls, and people routinely assume they are one
- Opting out of training stops future models learning from your chats
- It does not remove the chats. They sit in your account until you delete them
- Deletion then works through the back end over about a month, with exceptions
- At OpenAI there are four different things here: training, history, temporary chats, and deletion

Two Traps Worth Knowing

- Rating a Claude conversation with thumbs up or down sends the whole conversation into a five-year store, whatever your training setting says
 - Never rate a conversation that contains unpublished results
- ChatGPT's memory can hold more than the summary you are shown
 - Clearing it properly means chats, files, connected apps and the summary, separately

Text Is Starting to Be Watermarked

- A watermark is being added to Claude's output, to comply with European law
- It works by steering which of several equally good words gets picked, so nothing is added and nothing is visible
- It answers one question: did Claude probably generate this
- It fails on short passages, on factual writing, on code, and on anything rewritten
- So no watermark found is not evidence a person wrote it, and no error rate has been published

<https://www.anthropic.com/news/claude-text-watermark>

What Only One Side Has

- Only OpenAI: image generation, live two-way voice, speech and transcription, embeddings, and open-weight models you can run yourself
- Only Anthropic: Claude Science, which renders structures, submits jobs to a cluster, and runs an agent that flags citations and figures that do not match the code behind them
- Also only Anthropic: Claude Design, an org identity in Slack, and separate sessions that can pass findings to each other
- If you need images or speech that settles it, and the same goes for job submission and provenance

Two Advantages Being Withdrawn

- Video generation is ending. The app closed in April and the interface closes on 24 September, with no replacement
- Fine-tuning is closing: no new users now, and it stops entirely in January
- So if you were told you could train one of these on your own corpus, that is no longer true of either company's frontier models
- What is left instead: projects, caching, file search, packaged skills, or self-hosting an open model

Questions to Ask About Your Own Work

- Where do you fit in your own workflow? That is what sets the boundary of what should be automated
- What is actually in your account that you would not want in a training set?
- For each tool you would connect, what is the worst thing it could do with that access?
- What are your guardrails against a confident wrong answer?
- How do you check the output without reading every line of it?

Where to Go From Here

- Slides and materials: https://github.com/dlab-berkeley/AI_Pulse
- Workshop suggestions: <https://forms.gle/x8KfejgCViE7t6iQA>
- What Berkeley already gives you:
<https://life.berkeley.edu/ai-tools-for-berkeley-students>
- D-Lab: <https://dlab.berkeley.edu>

Next: AI in the News, 29 September. Then Teaching, Learning and Collaborating, 13 October.
Office hours next week, same day and time, on Zoom.

References

- Plans, models and prices: claude.com/pricing, platform.claude.com/docs, chatgpt.com/pricing and OpenAI's API pricing page. All checked September 2026
- Academic programs: OpenAI, ChatGPT for Academic Researchers, launched 29 July 2026; Anthropic, AI for Science and Team Plan for Scientists, claude.com/programs/team-plan-for-scientists
- Data and training defaults: both vendors' privacy and data-usage pages, checked September 2026
- Discontinuations: Sora API ends 24 September 2026; fine-tuning job creation ends 6 January 2027; Atlas closed 9 August 2026
- Berkeley's own tools: life.berkeley.edu/ai-tools-for-berkeley-students
- Auto mode and the 89 percent figure: Anthropic, "Auto mode is now the default in Claude Code", 7 August 2026. Red teaming by Apollo Research
- Permission modes, loops and project retrieval: code.claude.com/docs, support.claude.com
- Memory and watermarking: Anthropic, 25 August and 14 August 2026. The watermark uses Google DeepMind's SynthID-Text