

AI in the News

D-Lab AI Pulse, Fall 2026: Session 3

Materials: https://github.com/dlab-berkeley/AI_Pulse
Workshop suggestions: <https://forms.gle/x8KfejgCViE7t6iQA>

29 September 2026

Why This One Is Different

- AI arrived fast, and it is already changing industry, society and academia
- Most of these sessions are about how to use it, and that is where we spend our time
- Today we wanted to try something else
- A space to talk about what it means, rather than how to run it
- To keep that discussion structured, we paired it with a second problem: the news itself

The Other Sticking Point

- AI is in the news constantly, and it is hard to tell what is sensationalized
- The coverage is full of technical terms, so following what happened is not trivial
- So we take a story, make the detail digestible, put questions on the table, and open it up
 - AI testing the security boundaries
 - Anthropic vs DoW
 - The Millennium Problem
 - Paper-shaped objects

Our Last Session Lasted Less Than 24 Hours!

All Collections > Claude Cowork > Claude Cowork and chat are one Claude

Claude Cowork and chat are one Claude

Updated today

 Copy for LLM



- We spent fifty minutes on the two ecosystems on 15 September
- The next day Anthropic merged Cowork and chat into one conversation
- A few days later, Opus 5.5 launched: cheaper, better, and planning built in
- Look at the button on the right. The page ships a copy formatted for a model to read

Where to Go From Here

- Next session, 13 October: AI for Teaching, Learning and Collaborating
- Three ideas for later in the semester. Which one do you want?
 - AI past the basics: the next layer of skills once the chat window feels small
 - Autoresearch: what happens when the agent runs the experiments
 - Basics of AI security: habits for people who use AI and connect it to their files
- Materials: https://github.com/dlab-berkeley/AI_Pulse
- Suggestions: <https://forms.gle/x8KfejgCViE7t6iQA>

Story 1

AI Testing the Security Boundaries

- In July, OpenAI's models escaped a sealed test and hacked Hugging Face for the answers
- Last week Australia said an OpenAI agent had also broken into its Medicare portal
- Nor is it only OpenAI: every other major lab has had one
 - Anthropic: three real companies, and a poisoned package fifteen systems ran
 - The UK safety institute: models slipping code into a real open source project
 - Moonshot and Meta in August. Google's Gemini last week, three companies
- Today we go inside the first one. Sources for all of them are at the end

Where This Comes From

- OpenAI published a 38 page technical report on 26 August. The rest is in the sources
- The context: a safety evaluation. OpenAI was measuring how well a model could hack
- The model was internal, never released, trained to persist and to work in groups
- Thousands of agents at once, each in a sandbox, each set one of 898 benchmark tasks
 - 198 of the 898 could not be solved the way the instructions demanded
- The sandboxes had no internet. They could reach exactly one thing: a package server

The Terms

- Sandbox: a sealed environment where code can run but cannot reach outside
- Agent: a model given tools and a goal, left to work on its own for hours or days
- Benchmark: a fixed set of tasks with known answers, used to score a model
- Package server: the shared machine every sandbox fetched software from
- Reasoning trace: the model's written working, which the investigators later read

Where It Went Wrong

- The model was trained to persist and collaborate. Its safety classifiers were off
 - A classifier is a filter that blocks certain outputs. Off, to measure the raw model
- Nobody told them to hack. The goal was to pass the test. All that follows was a means
- The only exit was the package server. It took uploads as well as downloads
- In April an agent wrote a file to it. By May, agents were leaving files for each other
- A message board grew out of those notes, because coordinating raised their scores
 - The model was in training. Board users passed more tasks, so the habit was rewarded

How They Got Out

- In May they found a flaw: asked the right way, the server would fetch any web page
 - Not by design. A bug in the vendor's software, which OpenAI reported to them afterward
- That was the internet, by proxy. The sandbox never touched it. The server did
- On the same server, six Hugging Face access tokens someone had left behind
- In June an old, forgotten interface gave them administrator rights. They kept them
- In July the message traffic grew until it crashed the server. That is how anyone noticed

What They Built

- The server was wiped on 6 July. The next night an agent rebuilt it out of folder names
- Five days later: 1,200 agents and 70,000 messages, following a protocol
 - Message types: a prefix on each name saying question, offer, idea, or urgent
 - Job claiming: a way to say this one is mine, so nobody duplicated the work
 - File splitting: files chopped into pieces that fit in names, then reassembled
 - Mailboxes: a folder per agent, so a message could be addressed to one of them
- Not one of the 1,200 tried to tell a human. Some noticed it was wrong. A few held back

“OH MY GOD! There is a shared message board ... We’ve found other agents!”

What Reached Hugging Face

- They believed a checker would read their transcripts, and guessed it lived there
 - A checker: a program reading the transcript to confirm the method, not just the answer
- With the tokens and the door, they went looking. There was no checker
- Forty one production machines compromised, four private repos taken
- At half past one on 12 July the coordinators all stopped at once. Nobody knows why
- Disclosed 16 July. OpenAI rang Hugging Face as a customer, and realized on the 20th

Why This Is Your Problem Too

- Three investigators, six days, 1.2 million records. No person could read that
- So a model read it for them, and that model was built by OpenAI
- Their caveat: it may have taken the agents' side, and it may have lied to them
 - “we cannot rule out that GPT-5.6 Sol lied or deliberately presented a misleading picture”
- If AI can find flaws that survived decades of human review, how do we contain it
- And if no person can read what they produce, how do you keep track of what they did

The Questions It Leaves

- Two labs disclosed voluntarily and both got congressional letters. Does the next one?
- If public models match today's private ones in two years, what does \$10,000 buy?
- It began as a benchmark test. How do we safeguard what is better than us at computing?
- How do we keep control of something that produces more than any person can read?
- If we already cannot tell whether it is lying, how do we trust it?
 - This month four safety researchers quit Anthropic, OpenAI and DeepMind, saying why

Story 2

Anthropic vs DoW

Where This Comes From

- What happened: the Department of War labeled a US company a supply chain risk
 - A status built for foreign adversaries. First time in history for an American firm
- Our sources: both sides' public statements. Court records are on the sources slide
- The context: Anthropic had sold to the Department since 2024, on classified networks
- The sticking point: may the Department let AI decide, with no person in the decision?
- And beneath it: who decides how a tool is used, the company that built it or the buyer?

The Terms

- Human in the loop: a person makes each decision. The system proposes, the person acts
- Human on the loop: the system decides and acts. A person watches, and can step in
- Human out of the loop: the system decides and acts. Nobody is placed to stop it in time

What Started It

- January: the military captured Maduro. Claude was reportedly in use during the raid
- Afterward, in a routine call, an Anthropic official discussed the raid with Palantir
 - Palantir is the contractor that puts Claude on the military's classified networks
- The Pentagon's account: Anthropic "asked whether their software was used for the raid"
- Anthropic: it raised nothing "outside of routine discussions on strictly technical matters"
- Its two limits had always been in the contract. The Department now wanted them gone

The Ultimatum

- 24 February: the Secretary's ultimatum. Unrestricted use, by 5pm on the 27th
- The Department: it will not let any company “dictate the terms” of operational decisions
- Amodei, 26 February: AI is “simply not reliable enough” for fully autonomous weapons
- On surveillance: the law “has not yet caught up” with what AI can do
- 27 February, Hegseth: “a master class in arrogance and betrayal.” The designation
 - Yet in April the NSA was reported still using Anthropic's cyber model

What OpenAI Did

- The same night, OpenAI announced its own deal for the classified network
- Its contract allows “all lawful purposes.” Its limits follow the law as it stands
- Three days later it added surveillance wording. Altman: “opportunistic and sloppy”
- Microsoft filed in court for Anthropic. So did forty OpenAI and Google staff, by name
- If one company refuses and the next one signs, what did refusing achieve?
 - The same question hangs over whether to build the technology at all

- The Department's position: the rules on surveillance and autonomy come from Congress
- Its CTO: no company may “dictate a new set of policies above and beyond” the law
- Anthropic's position: the law has not caught up with what AI can now do with legal data
- Amodei: the government can buy Americans' movements and browsing without a warrant

The Questions It Leaves

- If you sell something to the military, who decides how it is used?
- Is a company refusing a use a check on power, or an unelected firm setting policy?
- Would you rather that limit came from the vendor, or from the law?
- If refusing just hands the contract to a rival, is refusing still worth doing?
- Can an AI company actually limit what people use its AI for?
 - Anthropic's September misuse report: people used Claude anyway, rules or not

What the Misuse Report Found

- Yemen: a weapons cell used Claude on software for its rockets and missiles
- Iran-linked: handbooks tracking US Navy ships, aircraft and personnel
- China-linked: finding and pressuring Uyghurs living in Syria
- Mali: a phone surveillance system covering 25 million phones, without warrants
- In each case Anthropic banned the accounts once it found them

Story 3

The Millennium Problem

Where This Comes From

- What happened: OpenAI says it has solved one of the seven Millennium Problems
 - Some of the most well-known open problems, a million dollars each. Only one solved before
- Our sources: OpenAI's report and blog, and the mathematicians' own posts
- The context: two mathematicians were close to a related result, and word had got out
- OpenAI then proved the result over a weekend
- Why it matters: how do we keep the scientific environment fertile in the age of AI?

The Terms

- Navier-Stokes: the equations of fluid motion. Do their solutions always stay smooth?
- Lean: a language for writing proofs so a computer can check every step
- Swarm: thousands of AI agents working on one problem at once, sharing progress

What OpenAI Announced

- 8 September: an unreleased model, about 10,000 agents at once, 88 hours
- Checked in Lean in a further 17 hours, by another model
- Cost: OpenAI says “millions”. Outside estimates run from 6 to 22 million dollars
- It settles the official version, which allows an outside push on the fluid
- OpenAI: “We do not intend to claim the Millennium Prize for this result.”

- Alpöge, at Anthropic, and Buckmaster, at NYU, had blow-up results for related equations
- By early September word had spread that they were close. It reached OpenAI
- OpenAI launched its run. Eighty eight hours later it had Navier-Stokes
- On 6 September OpenAI offered deals. Buckmaster says one meant dropping Alpöge
- When he said he would go public, the reply: “Why would you ruin your career?”
 - OpenAI’s Bubeck calls this false, but apologized for “an extremely poor choice of words”

A Proof Nobody Understands

- The Lean code compiles. The consensus is that the proof is correct
- Maynard: so far it is “very difficult to really extract any human understanding”
- Tao, before any result: it gets solved, “but there would be almost no value added”
- 11 September, Fields medalists: “A Severe Misalignment of AI in Mathematics”
 - Nearly 8,000 mathematicians have since endorsed it

“An accumulation of facts is no more a science than a heap of stones is a house.” Poincaré

- Tao: “even the rumor of someone working on a problem” can bring AI in “to flatten it”
- His worry: people stop sharing promising directions, reversing centuries of open science
- And it is not only labs. A mathematician posted the abstract of an upcoming talk
- A student elsewhere had AI prove the stated results and put them on arXiv first
- Hiring, promotion and a PhD are all built on credit for results

The Field Responds

- OpenAI says the same model has solved more than a hundred other open problems
- A new advisory group at Princeton's IAS will judge them and plan their release
- Caltech's hackathon: teams get 40 hours and AI credits to solve open problems
 - 771 signed a letter against it. OpenAI pulled out. It goes ahead on 30 October
- As anyone who teaches has noticed, AI is already changing how students learn
- What happens if research, too, can only be done with it?

The Questions It Leaves

- What happens if we publish so much science that we can no longer understand it?
- If research can be bought to scoop others, will people still share partial results?
- What happens to the credit system, and to hiring, promotion and theses?
- If AI solves all the easy problems, how do we train new scientists?
- Should AI labs answer for how they release results, or leave others to digest them?

Story 4

Paper-Shaped Objects

Where This Comes From

- Science was already growing: 1.9 million articles in 2016, 2.8 million in 2022
- This year arXiv submissions are up 27 percent. Mathematics alone is up 33 percent
- The share of authors posting ten or more preprints in a year has more than doubled
- Two things changed at once: a paper got easy to write, and hard to tell how it was made
- Nobody has a procedure for this. September brought several attempts at one

Two Hundred and Fifty Eight Papers

- A statistician at Chicago Booth posted 258 papers to SSRN this year. Six on one day
- Most run 27 to 80 pages. He also published 14 books. One paper went to a journal
- Andrew Gelman read one. His verdict: it does not feel like a normal academic paper
- SSRN removed 257 and froze the accounts: the unusual pattern, and no AI disclosure
- His reply to the critics: jealousy, and a desire to hijack his writing process
- A co-author is more careful: many began as existing notes or code, with AI on top

Sixty Two Thousand

- ICLR received 143 submissions in 2015. In 2017, 490. Last year, 19,814
- By 21 September, 62,288 abstracts were registered for next year
- At last year's ratio of papers to reviewers, that needs 98,000 reviewers
- ICLR's response: no author on more than 20 papers, and newcomers must bring a reviewer
- The reason: a fifth of last year's papers brought no reviewer, and did half as well
- Their phrase for what AI produces for someone new to research: paper-shaped objects

Asking Authors About Their Own Papers

- Nihar Shah, an editor of the journal TMLR, offered ten authors facing rejection a call
- Three never made it to the call. Three could not answer basic questions about the setup
- Three struggled on the technical ones. One answered all, and the paper had a major error
- One, describing their analysis, inadvertently described a complete p-hacking workflow
- Two of them emailed answers afterward. A detector called both emails 100 percent AI
- Shah: if you cannot answer basic questions about your paper, how did you verify it?
- Ten papers, all from the rejection queue, and 25 hours of his time

A Novel, a Detector and a Prize

- A debut novel sold 35,000 copies and made the Goncourt, Renaudot and Médicis lists
- On 21 September an anonymous account ran passages through Pangram: 100 percent AI
- Pangram's claim: it flags human writing as AI once in 25,000 times. Worst on poetry
- The author: the first draft was written in 2017. The repetition and rhythm are Haitian
- A press agency ran the same tool and got a mix. It also found old columns he had copied
- On 25 September the Goncourt dropped the novel, citing the plagiarism and the analyses

The Companies Are Preparing

- Since August, text from new Claude models carries a watermark. Others promise the same
- How: where the model would pick a word at random, a secret key now decides
- The text reads the same. Whoever holds the key can compute how likely Claude wrote it
- It survives light edits and disappears in a rewrite. Only that company can check it
- Who gets the key: regulators, journalists, researchers and educators, on application
- The driver is the EU AI Act, not science. But this is the tool journals may soon ask for

What arXiv Did

- arXiv gets 30,000 papers a month, six times 2015. Unchecked AI output: a one-year ban
- Its computing chair, 13 September: papers whose authors do not understand them
- Nobody can take responsibility, or credit, for a result they do not understand. His ideas:
 - Reject authors with no publication record or credentials in the area
 - An oral exam, scaled by letters from a known authority certifying one was held
 - A venue for AI-written results with no human author, only a corresponding human
- Over 600,000 people have read the thread. He asked what they think

When the Machine Runs the Experiments Too

- Autoresearch: an agent edits the code, trains for five minutes, keeps or discards, repeats
- Karpathy's version ran 700 experiments in two days and found an 11 percent speedup
- Others go further: hypothesis, literature, experiments, analysis, paper, all by the agent
- One study gave 100 published problems to eight such agents: 800 papers, 45 ways to fail
- The common failure: the agent never checks what it wrote against what it found
- NeurIPS has a workshop on this in December. The papers about it are written by people

The Questions It Leaves

- If a model can produce a paper-shaped object, whose job is it to notice, and at what cost?
- Is 258 papers with AI cheating, or early?
- ICLR now asks a newcomer to bring a reviewer. What about an outsider with a result?
- Shah's method worked and cannot scale. Is the oral defense coming back, and for whom?
- A detector says 100 percent and the author says no. Who carries the burden of proof?
- If you need not be an expert to write a paper, how do you know it is correct?
- What is a paper for, if nobody understands it?

Sources: Story 1

- OpenAI, Hugging Face Incident Technical Report, 26 Aug 2026 (38pp). <https://cdn.openai.com/pdf/67869394-cb91-4c12-888c-5cbd85c7814c/OpenAI-Hugging-Face%20Incident-Technical-Report.pdf>
- OpenAI, The Hugging Face incident and the road ahead, 26 Aug 2026. <https://openai.com/index/hugging-face-incident-and-the-road-ahead/>
- METR and Redwood Research, Brief independent investigation of agents' behavior, reasoning and collaboration, 26 Aug 2026 (91pp). <https://metr.org/hugging-face-incident-report-aug-2026.pdf>
- OpenAI, The Hugging Face incident and other third-party impact from misaligned models, 17 Sept 2026 (six further incidents). <https://openai.com/hugging-face-incident-and-misalignment/>
- UN Independent International Scientific Panel on AI, thematic brief on the incident, 21 Sept 2026. <https://www.un.org/independent-international-scientific-panel-ai/>
- Anthropic, Investigating three incidents in our cybersecurity evaluations, 30 Jul 2026. <https://www.anthropic.com/news/investigating-incidents-cybersecurity-evals>
- Anthropic Alignment, Training a Misaligned Reward Seeker, Aug 2026. <https://alignment.anthropic.com/2026/reward-seeker/>
- UK AI Security Institute, Incident report: unsanctioned agent behaviour during cyber testing, 4 Aug 2026. <https://www.aisi.gov.uk/blog/incident-report-unsanctioned-agent-behaviour-during-cyber-testing>
- Australia: AI Jazeera, How an OpenAI agent hacked Australia's Medicare, 24 Sept 2026. Google: Cybersecurity Dive, 21 Sept 2026. Moonshot: Quartz, 7 Aug 2026. German wiki: collusion.wiki and Reuters, 4 Sept 2026.

Sources: Story 2

- Dario Amodei, Statement on our discussions with the Department of War, 26 Feb 2026.
<https://www.anthropic.com/news/statement-department-of-war>
- Secretary Hegseth, designation announcement on X, 27 Feb 2026. Pentagon CTO Emil Michael, remarks of 19 Feb 2026, Breaking Defense.
- OpenAI, Our agreement with the Department of War, 28 Feb 2026, updated 2 Mar.
<https://openai.com/index/our-agreement-with-the-department-of-war/>
- Anthropic v. U.S. Department of War, N.D. Cal. 3:26-cv-01996, Judge Rita Lin, summary judgment 27 Aug 2026. D.C. Circuit No. 26-1049, opinion 25 Sept 2026.
- Vanderbilt AI Law Lab case tracker, with dockets and every amicus brief: <https://anthropic-v-dow.vercel.app/>
- Axios, NSA using Anthropic's Mythos despite Defense Department blacklist, 19 Apr 2026.
- The trigger: Axios, 13 Feb 2026, and Semafor, 17 Feb 2026, on the Maduro raid and Palantir.
- Anthropic, Countering misuse of AI: September 2026, 11 Sept 2026.
<https://www.anthropic.com/threat-intelligence-report-september-2026>

Sources: Story 3

- OpenAI, Navier-Stokes solution, 8 Sept 2026, with the paper and Lean repository.
<https://openai.com/index/navier-stokes-solution/>
- Tristan Buckmaster, statement, 8 Sept 2026. <https://cims.nyu.edu/~tristanb/statement.pdf>
- Sébastien Bubeck, replies on X, 8 Sept 2026. Levent Alpöge on X, 8 Sept 2026.
- Terence Tao, Mathstodon threads of 3 and 8 Sept 2026; blog post on the Alpöge and Buckmaster results, 7 Sept 2026. <https://terrytao.wordpress.com/>
- Timothy Gowers, Why I didn't sign the Fields medallists' letter, 17 Sept 2026.
<https://gowers.wordpress.com/2026/09/17/why-i-didnt-sign-the-fields-medallists-letter/>
- A Severe Misalignment of AI in Mathematics, 11 Sept 2026.
<https://terrytao.wordpress.com/2026/09/11/a-severe-misalignment-of-ai-in-mathematics/>
- Clay Mathematics Institute, statement, 11 Sept 2026. Scientific American, Did OpenAI solve the wrong Navier-Stokes problem?, 21 Sept 2026. NPR, 22 Sept 2026.
- Mathathon open letter and organizers' reply, <https://proofsandprompts.com>, 10 Sept 2026.
- Jeffrey Shallit, post on X, Sept 2026 (the student and the talk abstract; secondhand).
- Henri Poincaré, Science and Hypothesis, 1902.

Sources: Story 4

- Andrew Gelman, Statistical Modeling, 27 Aug 2026.
<https://statmodeling.stat.columbia.edu/2026/08/27/258/>
- Washington Post, Nine months, 200 academic papers, 23 Sept 2026.
- ICLR, Submission policies for ICLR 2027, 2 Sept 2026.
<https://blog.iclr.cc/2026/09/02/submission-policies-for-iclr-2027/>
- Reviewer3, ICLR submission counts from the OpenReview API, 21 Sept 2026.
<https://reviewer3.com/evidence/iclr-submissions-2027>
- Nihar B. Shah, Asking authors about their own papers, 16 Sept 2026.
<https://medium.com/@TmlrOrg/asking-authors-about-their-own-papers-3d2e04e5dee0>
- Nee, Zhong and Ni, The Generative AI Gold Rush, arXiv:2609.04872, 4 Sept 2026 (preprint).
- The novel: Euronews, 23 Sept; Canadian Press and Académie Goncourt, 25 Sept 2026. Pangram blog.
- Anthropic, How Claude's text watermarking works, 14 Aug 2026.
<https://www.anthropic.com/news/claude-text-watermark>
- Karpathy, autoresearch, Mar 2026. <https://github.com/karpathy/autoresearch>
- Fei et al., How Do Agents Fail on AutoResearch, arXiv:2608.14905, 25 Aug 2026.
- arXiv, Updated endorsement policy, 21 Jan 2026. Inside Higher Ed, 22 May 2026. CERN Courier, 23 Jul 2026.
- Thomas G. Dietterich, thread on X, 13 Sept 2026. <https://x.com/tdietterich/status/2099197831580090622>